

ST1516 – DevOps & AI Automation
Jeyakumar Sriram
DAAA/FT/2B/01
2214618

Giggles Gears

Where you can find you future car...





Table of contents

01

Application

02

DevOps Process

03

Regression

04

Flask

05

Testing

06

Advanced




An illustration of a light-skinned hand holding a dark blue car key with a silver blade. The hand is positioned on the left side of the frame.

01

An illustration of a dark-skinned hand reaching out from the right side of the frame, palm facing up.

Application

The Dataset

 ADITYA · UPDATED 3 YEARS AGO

801New NotebookDownload (1 MB)

100,000 UK Used Car Data set

100,000 scraped used car listings, cleaned and split into car make.

[Data Card](#)[Code \(296\)](#)[Discussion \(14\)](#)

About Dataset

If you download/use the data set I'd appreciate an up vote, cheers.

Updated

Scraped data of used cars listings. 100,000 listings, which have been separated into files corresponding to each car manufacturer. I collected the data to make a tool to predict how much my friend should sell his old car for compared to other stuff on the market, and then just extended the data set. Then made a more general car value regression model.

previous version

Picked two fairly common cars on the British market for analysis (Ford Focus and Mercedes C Class). The hope is to find info such as: when is the ideal time to sell certain cars (i.e. at what age and mileage are there significant drops in resale value). Also can make comparisons between the two, and make a classifier for a ford or Mercedes car. Can easily add more makes and models, so comment for any request e.g. if you want a big data set of all Mercedes makes and models.

Content

The cleaned data set contains information of price, transmission, mileage, fuel type, road tax, miles per gallon (mpg), and engine size. I've removed duplicate listings and cleaned the columns, but have included a notebook showing the process and the original data for anyone who wants to check/improve my work.

Usability 10.00

License
CC0: Public Domain

Expected update frequency
Quarterly

Tags
Automobiles and Vehicles

Inspiration

This dataset consists of prices and other related information on a 100000 cars from the UK. Since this dataset is from the UK most of the brands are from there.

I had to join and clean them to properly format to one df

	model	year	price	transmission	mileage	fuelType	tax	mpg	engineSize	brand
0	A1	2017	12500	Manual	15735	Petrol	150	55.4	1.4	audi
1	A6	2016	16500	Automatic	36203	Diesel	20	64.2	2.0	audi
2	A1	2016	11000	Manual	29946	Petrol	30	55.4	1.4	audi
3	A4	2017	16800	Automatic	25952	Diesel	145	67.3	2.0	audi
4	A3	2019	17300	Manual	1998	Petrol	145	49.6	1.0	audi

Application Demo

The application I built has the following key features.

Users can:

- browse a catalogue of cars available for resale
- add their desired cars to a wishlist
- get a rough estimate of how much a car with their desired features would cost
- get 3 recommended cars based on their input on desired features
- view and delete their prediction history
- view and delete items from wishlist
- sign up or login to save their progress
- view and modify their profile




<https://giggles-gears.onrender.com/>

Some pictures

Giggles Gears

[Home](#) [Catalogue](#) [SmartSearch](#) [History](#) [Wishlist](#) [Profile](#) [Login/Signup](#)



Want to know how much your dream car costs?!

We've got you covered!

[Take the quiz!](#)

[Login/Signup](#)

To save your desired cars, please login or sign up.

Giggles Gears

[Home](#) [Catalogue](#) [SmartSearch](#) [History](#) [Wishlist](#) [Profile](#) [Login/Signup](#)

Our Cars Collection

Car Name

Car Brand

Fuel Type

Taxable?

Min Value

Max Value

search

Results

audi A1

6 years old

15735 km Mileage

55.4 MPG

£12500.0

Add to Wishlist

audi A6

7 years old

36203 km Mileage

64.2 MPG

£16500.0

Add to Wishlist

audi A1

7 years old

29946 km Mileage

55.4 MPG

£11000.0

Add to Wishlist

audi A4

6 years old

25952 km Mileage

67.3 MPG

£16800.0

Add to Wishlist

audi A3

4 years old

1998 km Mileage

49.6 MPG

£17300.0

Add to Wishlist

audi A1

7 years old

32260 km Mileage

58.9 MPG

£13900.0

Add to Wishlist

audi A6

7 years old

76788 km Mileage

Add to Wishlist

audi A4

7 years old

75185 km Mileage

Add to Wishlist

audi A3

8 years old

46112 km Mileage

Add to Wishlist

History of Predictions										
Date	Year	Transmission	Mileage	Fuel Type	Tax	MPG	Engine Size	Brand	Price	Remove?
2023-12-03 13:22:34.793995	1	Manual	1	Petrol	1	1	1.0	audi	27882.9	Remove
2023-12-03 16:15:52.419472	3	Manual	35000	Petrol	150	30	0.5	audi	20346.51	Remove

Your Wishlist					
Car Name	Age	Brand	Price	Added On	Remove?
A1	7	audi	11000.0	2023-12-03 13:20:13.257672	Remove
A1	6	audi	12995.0	2023-12-03 13:22:36.983410	Remove

Giggles Gears

[Home](#) [Catalogue](#) [SmartSearch](#) [History](#) [Wishlist](#) [Profile](#) [Login/Signup](#)

prediction

Car Age

3

Type of Transmission

Manual

Mileage

35000

Engine Fuel

Petrol

Road Tax (£)

150

Miles Per Gallon

30

Engine Size in Litres

0.5

Car Brand

Audi

[Predict](#)

The price of a car with your specification will be around...
£ 20,346.51

Here are our recommendations:

audi A1

6 years old

24386 km Mileage

67.3 MPG

£12995.0

Add to Wishlist

audi A1

6 years old

22488 km Mileage

67.3 MPG

£12990.0

Add to Wishlist

audi A1

6 years old

26950 km Mileage

67.3 MPG

£12950.0

Add to Wishlist


An illustration featuring two hands, one light orange and one brown, exchanging a dark blue car key. In the center, a white square box with an orange border contains the number '02' in orange. The background is a light beige color with stylized white clouds, small orange four-pointed stars, and two orange trees on a rolling orange hill at the bottom. The entire scene is framed by orange borders on the left and right sides.


02


DevOps Process


Git & Scrum


Branches


releases  `bc933bd5` · Merge branch 'main' into releases · 1 hour ago


main  `default` `protected`
`0f5b7d41` · edited .gitlab-ci.yml file · 1 hour ago

deployment  `598207ff` · renamed it lol · 1 hour ago


tests  `3ffdc76b` · finished range testing · 2 hours ago


backend  `9123acdc` · recommend cars working · 13 hours ago

regressionML  `337b0ec1` · updated notebook with cos similarity · 14 hours ago

userauth  `1d01f931` · working auth and history save · 1 day ago

Releases branch has fully working user ready products. Main is more like a development branch. I'm using the centralized workflow from notes

database  `c7a90e69` · added database tables · 2 days ago

front-end  `a20b73c1` · added showcase · 2 days ago

Product Enhancement
Nov 30, 2023–Dec 4, 2023
`Closed` 4618-devops / CA1-2B01-2214618-JeyakumarSriram










































Working User System
Nov 24, 2023–Dec 2, 2023
`Closed` 4618-devops / CA1-2B01-2214618-JeyakumarSriram










































Create new base application
Nov 17, 2023–Nov 30, 2023
`Closed` 4618-devops / CA1-2B01-2214618-JeyakumarSriram










































Minimum Viable Product
Nov 4, 2023–Nov 11, 2023
`Closed` 4618-devops / CA1-2B01-2214618-JeyakumarSriram










































Issue List -> Scrum Board






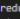





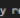
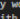




















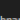







Development

Open                                         

To Do                                         

In Progress                                         

Testing                                         

Closed                                         

Each issue has a user story if you read the description. Each issue has a deadline and a sprint assigned. "Open" section is my product backlog. "To Do" is my sprint backlog

Milestones -> Sprints

Since I changed my dataset halfway I had to create a sprint to adapt my code to the new data

Git Commits

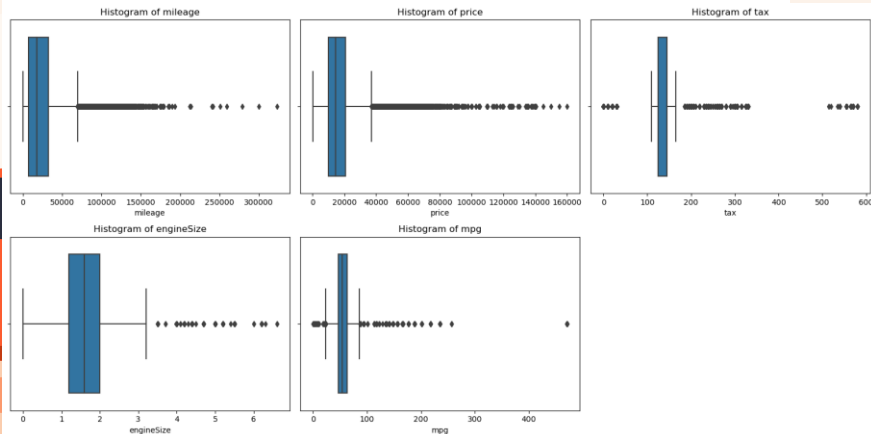


```
bc933bd (tag: v3.0, origin/releases, releases) Merge branch 'main' into releases
|\
| * 0f5b7d4 (HEAD -> main, origin/main, origin/HEAD) edited .gitlab-ci.yml file
| * b223c77 Updated .gitlab-ci.yml file
| * 598207f (origin/deployment, deployment) renamed it lol
| * 5ef4afe added cicd
| * 3ffdc76 (origin/tests, tests) finished range testing
| * ad24eb7 done api tests
| * 69f9988 add user working and fixed bugs
| * 6144a66 finished consistency testing
| * 310f994 expected and unexpected works
| * 3164f3a removed redundant
| * d95acc7 database testing works
| * 4046e83 (tag: v2.0) Merge branch 'main' into releases
|\
| * c964459 refactored code and formatted
| * 549b2f7 Merge branch 'backend'
|\
| * 9123acd (origin/backend, backend) recommend cars working
| * 5907fa3 added wishlist features
| * 779b02a history removal works
| * 44abaf7 history works
| * 5142908 done with wishlist
| * 444d1fc finished profile and update
| * 4eb926c preloading cars done
| * 337b0ec (origin/regressionML, regressionML) updated notebook with cos similarity
|/
| * a0283ab Merge branch 'deployment'
|\
| * 4a04d0d unicorn server works
| * 6f76f43 decide tests and created apis
| * 3ecb8b6 slight refactoring
| * 27fcae1 refactored code for testing
| * 1c1c765 Merge branch 'main' into tests
|\
|/
| * 1d01f93 (origin/userauth, userauth) working auth and history save
| * 0e53537 Merge branch 'database'
|\
|/
| * c7a90e6 (origin/database, database) added database tables
| * a20b73c (origin/front-end, front-end) added showcase
| * bd658d2 added wishlist
| * a51a286 added history page
| * d8bfa32 formatted prediction page
| * 2accff4 added homepage
| * 307ce59 added nav and footer
|/
|/
| * 65c7a1e init tests
|/
|/
| * 25e0f61 (tag: v1.0) Added Readme
| * ff2ea0d fixed issues with model and integrated
| * 9cb8aef Merge branch 'regressionML' into 'main'
|\
| * df9d7c9 Merge branch 'main' into 'regressionML'
|/
|/
| * 2b52804 Merge branch 'wireFrames' into 'main'
|\
|/
| * 75fc283 added sign up
| * b4785f3 added login
| * aac5a6f Added signup.drawio
```


An illustration featuring two hands. The left hand, in a light orange tone, holds a dark blue car key. The right hand, in a brown tone, is reaching out towards the key. In the center, a white square box with an orange border contains the number '03' in orange. The background is a light beige color with stylized white clouds, small orange stars, and a landscape at the bottom with rolling hills and two orange trees.

03


Regression

Exploratory Data Analysis



Based on the EDA:

- Some classes had very low rows and decided to remove them.
- Outliers in year and other numeric columns
- Price highly correlated with year and engine size
- No null values
- Model has too many classes and cannot be used



Preprocessing



Cleaning

- Removed outliers
- Remove classes with low representation
- Removed models with very little rows

```
df = df[~df["fuelType"].isin(["Other", "Electric"])]
df = df[~df["transmission"].isin(["Other"])]
df = df[df["year"] > 1990]
df = df[df["year"] < 2023]
df["year"] = 2023 - df["year"]
df = df[df["engineSize"] != 0]
df = df[df["mpg"] < 300]
df = df[
    ~df["model"].isin(
        df["model"].value_counts()[df["model"].value_counts() < 200].index
    )
]
```




Splitting

- Used train-val-test split with ratio 70-10-10
- Dropped model column from dataset

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, train_size=0.70, random_state=99
)

X_test, X_val, y_test, y_val = train_test_split(
    X_test, y_test, train_size=0.5, random_state=99
)

features = X_train.columns
```



Pipeline


- Performed ordinal encoding for transmission
- Performed one hot for others
- Scaled numeric columns

```
ord_enc = ["transmission"]
onehot_enc = ["brand", "fuelType"]

scale_trans = ["mileage", "tax", "mpg", "engineSize"]

preprocessor = ColumnTransformer(
    transformers=[
        (
            "label_enc",
            OrdinalEncoder(categories=[["Manual", "Semi-Auto", "Automatic"]]),
            ord_enc,
        ),
        ("onehot_enc", OneHotEncoder(), onehot_enc),
        ("scale_trans", StandardScaler(), scale_trans),
    ],
    remainder="passthrough",
)
```

Modelling



Try many

- Compared r2 and MAPE from various models
- Random Forest, ExtraTrees and Gradient Boosting performed well on val

	name	r2	mape
0	AdaBoostRegressor	0.561832	0.379821
1	Elasticnet	0.720060	0.229378
2	ExtraTreesRegressor	0.945363	0.082342
3	GradientBoosting	0.901410	0.129191
4	Lasso	0.797757	0.212168
5	LinearRegression	0.797793	0.212352
6	RandomForest	0.949272	0.079728
7	Ridge	0.797782	0.212236



VotingRegressor

- Used three best models just now to make it
- Performed around the same
- But takes very long to train and predict

```
voting_regressor = VotingRegressor(  
    estimators=[  
        ("ExtraTreesRegressor", ExtraTreesRegressor()),  
        ("GradientBoosting", GradientBoostingRegressor()),  
        ("RandomForest", RandomForestRegressor()),  
    ]  
)  
vote_pipe = getPipe("voting_regressor", voting_regressor)  
vote_pipe.fit(X_train, y_train)  
  
y_val_pred = vote_pipe.predict(X_val)  
  
r2 = r2_score(y_val, y_val_pred)  
mape = mean_absolute_percentage_error(y_val, y_val_pred)  
  
print("Voting Regressor")  
print(f"R2: {r2}")  
print(f"MAPE: {mape}")  
  
Voting Regressor  
R2: 0.9488956222712608  
MAPE: 0.08500218878528085
```



LightGBM

- Used from external library
- Performed around the same
- But is noticeably faster

```
lgbm_regressor = LGBMRegressor()  
lgbm_pipe = getPipe("lgbm_regressor", lgbm_regressor)  
lgbm_pipe.fit(X_train, y_train)  
  
y_val_pred = lgbm_pipe.predict(X_val)  
  
r2 = r2_score(y_val, y_val_pred)  
mape = mean_absolute_percentage_error(y_val, y_val_pred)  
  
print("LightGBM Regressor")  
print(f"R2: {r2}")  
print(f"MAPE: {mape}")  
  
LightGBM Regressor  
R2: 0.9392643657067584  
MAPE: 0.09735651323184125
```

Random Forest vs LightGBM

Random Forest

- Performed slightly better than LGBM
- Slower than LGBM
- More accuracy might not be that significant to user than speed

```
rf_regressor = getPipe("Random Forest", RandomForestRegressor())
lgbm_regressor = getPipe("LightGBM", LGBMRegressor())

start_time_fit = time.time()
rf_regressor.fit(X_train, y_train)
fit_time_rf = time.time() - start_time_fit

start_time_fit = time.time()
lgbm_regressor.fit(X_train, y_train)
fit_time_lgbm = time.time() - start_time_fit

start_time_predict = time.time()
y_val_pred = rf_regressor.predict(X_val)
predict_time_rf = time.time() - start_time_predict

start_time_predict = time.time()
y_val_pred = lgbm_regressor.predict(X_val)
predict_time_lgbm = time.time() - start_time_predict

print(f"Fit Time: {fit_time_rf:.4f} seconds - Random Forest")
print(f"Fit Time: {fit_time_lgbm:.4f} seconds - LightGBM")

print(f"Predict Time: {predict_time_rf:.4f} seconds - Random Forest")
print(f"Predict Time: {predict_time_lgbm:.4f} seconds - LightGBM")
```

```
Fit Time: 37.1690 seconds - Random Forest
Fit Time: 0.5411 seconds - LightGBM
Predict Time: 0.7888 seconds - Random Forest
Predict Time: 0.0490 seconds - LightGBM
```

LightGBM

- Performed slightly worse than RF
- 40x faster at training
- 20x faster at prediction
- Will provide better user experience



Tuning & Final Model

- Since we making final model I will use the entire dataset without splits
- Use RandomisedCV for speed as the slight increase in accuracy with GridsearchCV may not be significant
- Best r2 is better than RF just now
- Final pipeline is ready and will be saved as pickle file to be loaded by Flask

```
merged_X_train = pd.concat([X_train, X_val], axis=0, ignore_index=True)
merged_y_train = np.concatenate((y_train, y_val))
```

```
lgbm_regressor = getPipe("LightGBM", LGBMRegressor())

param_dist = {
    "LightGBM__learning_rate": [0.05, 0.1, 0.2],
    "LightGBM__n_estimators": [50, 100, 150],
    "LightGBM__max_depth": [-1, 5, 10, 20],
    "LightGBM__colsample_bytree": [0.7, 0.8, 0.9],
    "LightGBM__subsample": [0.7, 0.8, 0.9],
    "LightGBM__min_child_samples": [1, 5, 10],
    "LightGBM__num_leaves": [20, 40, 60, 80, 100],
    "LightGBM__reg_alpha": [0, 0.01, 0.03],
}
```

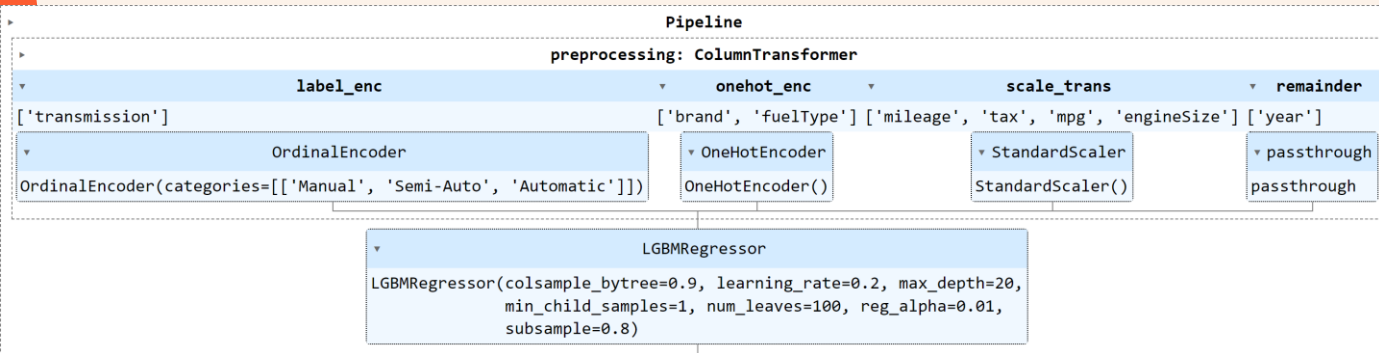
```
random_search = RandomizedSearchCV(
    estimator=lgbm_regressor,
    param_distributions=param_dist,
    scoring="r2",
    n_iter=10,
    cv=3,
    random_state=99,
)
```

```
random_search.fit(merged_X_train, merged_y_train)
```

```
best_params = random_search.best_params_
best_model = random_search.best_estimator_
```

```
y_pred = best_model.predict(X_test)
```

```
r2 = r2_score(y_test, y_pred)
print(f"Best Parameters: {best_params}")
print(f"R2 on Test Set: {r2}")
```



Best Parameters: {'LightGBM__subsample': 0.8, 'LightGBM__reg_alpha': 0.01, 'LightGBM__num_leaves': 100, 'LightGBM__n_estimators': 100, 'LightGBM__min_child_samples': 1, 'LightGBM__max_depth': 20, 'LightGBM__learning_rate': 0.2, 'LightGBM__colsample_bytree': 0.9}

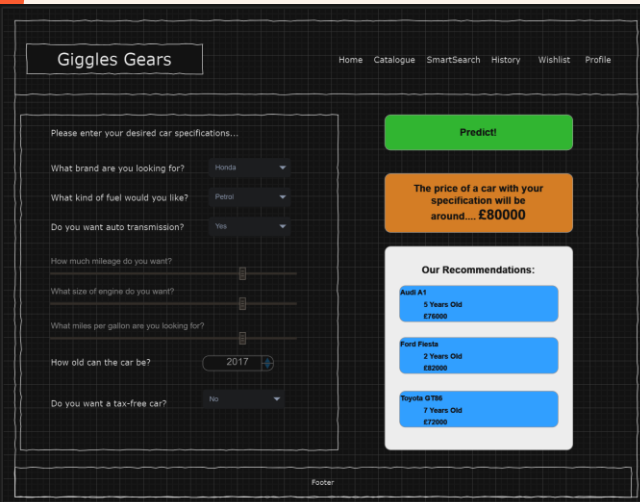
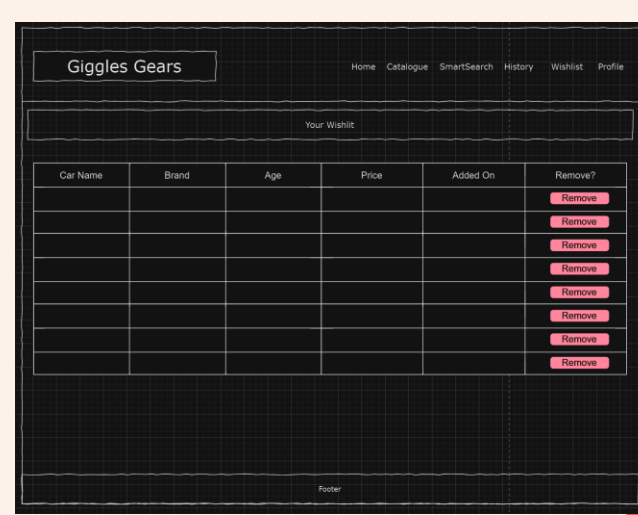
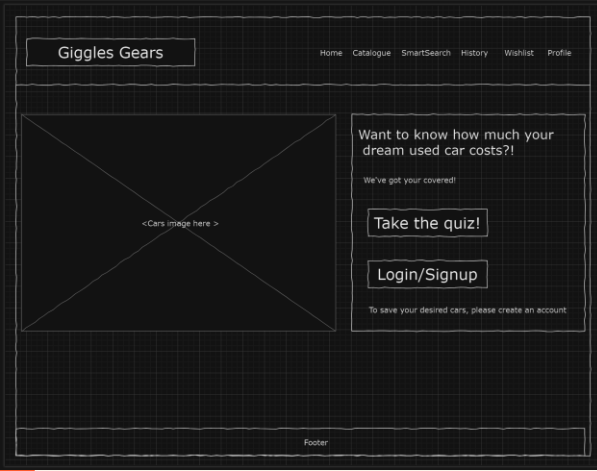
R2 on Test Set: 0.9438949546507968

An illustration featuring two hands. On the left, a light-skinned hand holds a dark blue car key. On the right, a darker-skinned hand is open, palm up, as if receiving the key. In the center, a white square with an orange border contains the number '04' in orange. Below this, the word 'Flask' is written in a bold, dark blue font. The background is a light beige color with stylized white clouds, small orange stars, and two orange trees on a rolling orange hill at the bottom. The entire scene is framed by orange borders on the left and right sides.

04

Flask

Wireframes



Flask App Setup

I used Flask-Login for user management and Flask-Bcrypt for hashing passwords

Application Structure

```
C:\
|
| api.py
| auth.py
| config.cfg
| database.py
| fileStruct.txt
| forms.py
| ml.py
| models.py
| routes.py
| __init__.py
|
+---static
|   | cars.jpg
|   |
|   \---css
|       nav.css
|
+---templates
|   | base.html
|   | history.html
|   | index.html
|   | login.html
|   | prediction.html
|   | profile.html
|   | showcase.html
|   | signup.html
|   | wishlist.html
|   |
|   \---includes
|       footer.html
|       nav.html
|
\---pycache
    api.cpython-311.pyc
    auth.cpython-311.pyc
    database.cpython-311.pyc
    forms.cpython-311.pyc
    ml.cpython-311.pyc
    models.cpython-311.pyc
    routes.cpython-311.pyc
    __init__.cpython-311.pyc
```

Routes

1. Home Page

- **Route:** /
- **Description:** Renders the home page of the application.

2. History Page

- **Route:** /history
- **Description:** Displays the user's historical car data if logged in; otherwise, redirects to the login page.

3. Wishlist Page

- **Route:** /wishlist
- **Description:** Displays the user's wishlist if logged in; otherwise, redirects to the login page.

4. Showcase Page

- **Route:** /showcase
- **Description:** Allows users to search and filter cars. Displays a catalog of cars based on the search criteria.

5. Prediction Page

- **Route:** /predict
- **Description:** Performs car price prediction using a regression model. If the user is logged in, adds the prediction to their history. Displays the prediction result and recommends cars.

6. Login Page

- **Route:** /login
- **Description:** Provides user authentication. Users can log in with their credentials.

7. Signup Page

- **Route:** /signup
- **Description:** Allows new users to sign up for the application.

8. Profile Page

- **Route:** /profile
- **Description:** Displays and allows users to update their profile information if logged in; otherwise, redirects to the login page.

Database Schema

User Table:

- **id:** Integer (Primary Key)
- **name:** String(255) (Not Null)
- **email:** String(255) (Unique, Not Null)
- **password:** String(255) (Not Null)

Car Table:

- **id:** Integer (Primary Key)
- **year:** Integer
- **model:** String(255)
- **transmission:** String(255)
- **mileage:** Integer
- **fuelType:** String(255)
- **tax:** Integer
- **mpg:** Integer
- **engineSize:** Float
- **brand:** String(255)
- **price:** Float

Wishlist Table:

- **user_id:** Integer (Foreign Key, References User.id, Primary Key, Not Null)
- **car_id:** Integer (Foreign Key, References Car.id, Primary Key, Not Null)
- **date:** DateTime (Default: Current UTC DateTime, Not Null)

History Table:

- **histID:** Integer (Primary Key, Not Null)
- **id:** Integer (Foreign Key, References User.id, Not Null)
- **year:** Integer
- **transmission:** String(255)
- **mileage:** Integer
- **fuelType:** String(255)
- **tax:** Integer
- **mpg:** Integer
- **engineSize:** Float
- **brand:** String(255)
- **price:** Float
- **date:** DateTime (Default: Current UTC DateTime, Not Null)

An illustration of a light-skinned hand holding a dark blue car key with a silver blade. The background is a light beige color with stylized white clouds and small orange starburst shapes.

05

An illustration of a dark-skinned hand reaching out towards the center. The background is a light beige color with stylized white clouds and small orange starburst shapes.

Testing

API Tests:

- **User:**
 - Test add user
 - Test login API
- **Prediction:**
 - Test the prediction API

Range Tests:

- **Prediction:**
 - Try inputting large values and negative values

Consistency Tests:

- **Prediction:**
 - Input the same model with the same data multiple times and check consistency
- **API:**
 - Test if search results are consistent no matter how many times you request

Expected Scenarios:

- **Prediction:**
 - Missing values for the ML model
- **Routes:**
 - Invalid request types like PUT and DELETE for routes
- **Forms:**
 - Try every form with missing data
- **Database:**
 - Init models with missing params

Unexpected Scenarios:

- **User:**
 - Add a user
 - Login and obtain user ID
- **Prediction:**
 - Predict fixed input output
- **Database:**
 - Test initialization for each database class
 - Insert into the database and verify values
 - Remove from the database and ensure consistency
 - Edit the database and validate changes
 - Test populating cars DB

I did validity testing too but they are spread out between the diff files cuz its hard to separate them

```
import pytest
from application import app as application
from application import db as database

@pytest.fixture
def app():
    yield application

@pytest.fixture
def client(app):
    return app.test_client()

@pytest.fixture
def db(app):
    with app.app_context():
        yield database.session
```

Tests (Overview)

I used simple fixtures in my application. I group all the tests in different files based on what they are testing for like range or consistency.

All the passable tests have passed and the rest xfailed. You can ignore the warning as they are future warning about a sklearn function getting deprecated



```
===== test session starts =====
platform win32 -- Python 3.11.5, pytest-7.4.3, pluggy-1.3.0
rootdir: C:\Users\justf\Desktop\Projects\github\2b01-2214618-jeayakumarsriram-ca1
collected 61 items

tests\test_api.py ... [ 4%]
tests\test_consistency.py .... [ 11%]
tests\test_expected.py xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx [ 81%]
tests\test_range.py xx [ 85%]
tests\test_unexpected.py ..... [100%]

===== warnings summary =====

tests/test_api.py: 4 warnings
tests/test_consistency.py: 42 warnings
tests/test_range.py: 6 warnings
tests/test_unexpected.py: 4 warnings
C:\Users\justf\Desktop\Projects\github\2b01-2214618-jeayakumarsriram-ca1\env\Lib\site-packages\sklearn\utils\validation.py:605: FutureWarning
: is_sparse is deprecated and will be removed in a future version. Check `isinstance(dtype, pd.SparseDtype)` instead.
  if is_sparse(pd_dtype):

tests/test_api.py: 4 warnings
tests/test_consistency.py: 42 warnings
tests/test_range.py: 6 warnings
tests/test_unexpected.py: 4 warnings
C:\Users\justf\Desktop\Projects\github\2b01-2214618-jeayakumarsriram-ca1\env\Lib\site-packages\sklearn\utils\validation.py:614: FutureWarning
: is_sparse is deprecated and will be removed in a future version. Check `isinstance(dtype, pd.SparseDtype)` instead.
  if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):

tests/test_api.py: 1 warning
tests/test_consistency.py: 24 warnings
tests/test_unexpected.py: 1 warning
C:\Users\justf\Desktop\Projects\github\2b01-2214618-jeayakumarsriram-ca1\env\Lib\site-packages\sklearn\utils\validation.py:767: FutureWarning
: is_sparse is deprecated and will be removed in a future version. Check `isinstance(dtype, pd.SparseDtype)` instead.
  if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():

-- Docs: https://docs.pytest.org/en/stable/how-to/capture-warnings.html
===== 16 passed, 45 xfailed, 138 warnings in 3.84s =====
```

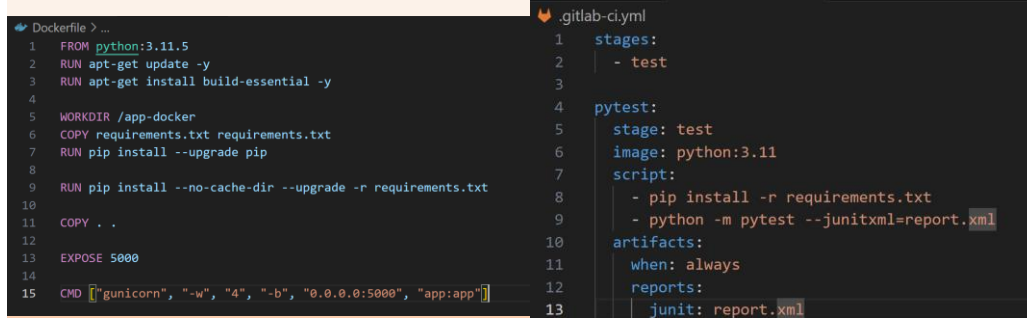
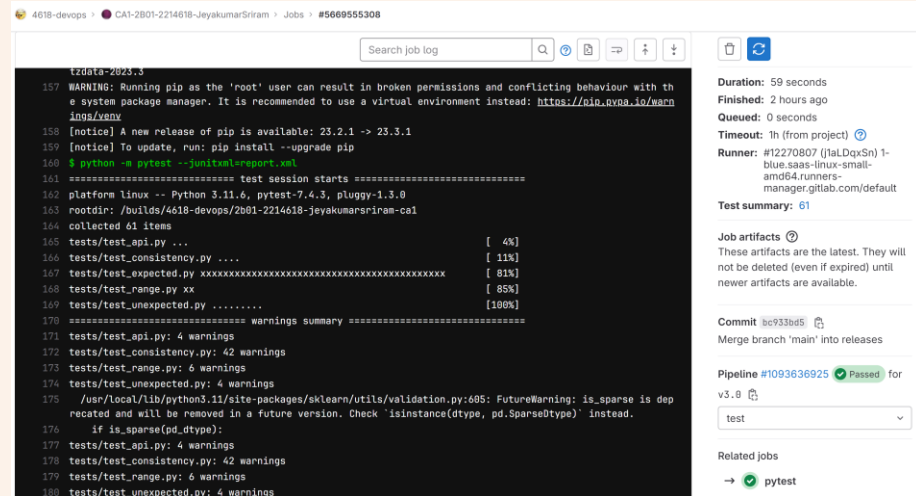
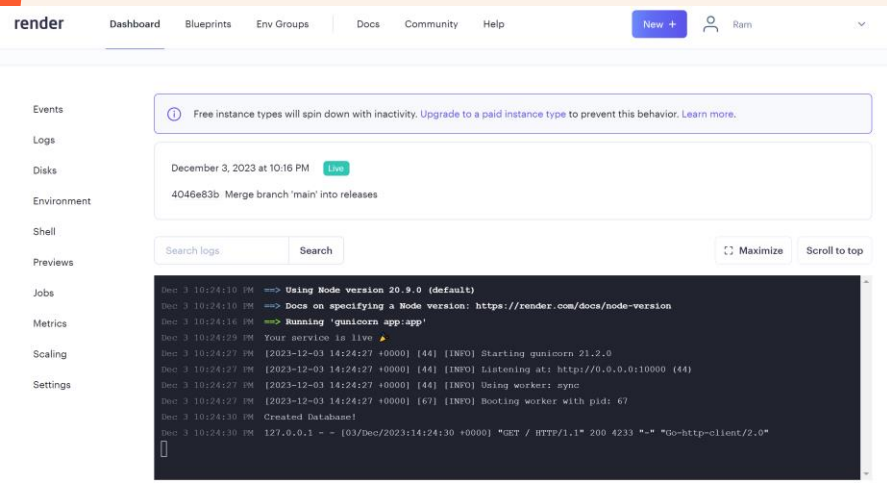


06

Advanced

Advanced

- Automatic Testing: I used gitlab's CI pipeline feature to construct one that has just one stage which is testing. Basically whenever I merge a branch or commit the tests I've created will be run and produce a report like this ==>>>
- Dockerise: throughout my development I used docker to avoid touching my local environment. Using unicorn
- Deployed to Render.com using GitLab. Initially I tried to deploy my image to dockerhub and then to render but to make things simpler I just connected my render to my releases branch to deploy every release



Thanks!

Do you have any questions?

youremail@freepik.com

+34 654 321 432

yourwebsite.com

CREDITS: This presentation template was created by **Slidesgo**, and includes icons by **Flaticon**, and infographics & images by **Freepik**

Please keep this slide for attribution

